



# Wine Alcohol Level Prediction

Yagiz Sezersan



# Objective

Predict alcohol level content from white wines given relevant data regarding the wine qualities and compositions.



# Data Overview

We get the data set from the [UCI Machine Learning Repository](https://archive.ics.uci.edu/ml/datasets/wine+quality) and the link to the data set is following:

<https://archive.ics.uci.edu/ml/datasets/wine+quality>

[Wine Quality - UC Irvine Machine Learning Repository](https://archive.ics.uci.edu/ml/datasets/wine+quality)


	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates	alcohol	quality
1	7.0	0.270	0.36	20.70	0.045	45.0	170.0	1.0010	3.00	0.45	8.8	6
2	6.3	0.300	0.34	1.60	0.049	14.0	132.0	0.9940	3.30	0.49	9.5	6
3	8.1	0.280	0.40	6.90	0.050	30.0	97.0	0.9951	3.26	0.44	10.1	6
4	7.2	0.230	0.32	8.50	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	6
5	7.2	0.230	0.32	8.50	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	6
6	8.1	0.280	0.40	6.90	0.050	30.0	97.0	0.9951	3.26	0.44	10.1	6
7	6.2	0.320	0.16	7.00	0.045	30.0	136.0	0.9949	3.18	0.47	9.6	6
8	7.0	0.270	0.36	20.70	0.045	45.0	170.0	1.0010	3.00	0.45	8.8	6
9	6.3	0.300	0.34	1.60	0.049	14.0	132.0	0.9940	3.30	0.49	9.5	6
10	8.1	0.220	0.43	1.50	0.044	28.0	129.0	0.9938	3.22	0.45	11.0	6
11	8.1	0.270	0.41	1.45	0.033	11.0	63.0	0.9908	2.99	0.56	12.0	5
12	8.6	0.230	0.40	4.20	0.035	17.0	109.0	0.9947	3.14	0.53	9.7	5
13	7.9	0.180	0.37	1.20	0.040	16.0	75.0	0.9920	3.18	0.63	10.8	5
14	6.6	0.160	0.40	1.50	0.044	48.0	143.0	0.9912	3.54	0.52	12.4	7
15	8.3	0.420	0.62	19.25	0.040	41.0	172.0	1.0002	2.98	0.67	9.7	5
16	6.6	0.170	0.38	1.50	0.032	28.0	112.0	0.9914	3.25	0.55	11.4	7
17	6.3	0.480	0.04	1.10	0.046	30.0	99.0	0.9928	3.24	0.36	9.6	6
18	6.2	0.660	0.48	1.20	0.029	29.0	75.0	0.9892	3.33	0.39	12.8	8

## Data Overview

The data set have total  
4898 rows and 12 columns  
as shows in the figure

```
> nRows <- nrow(data6)
> print(nRows)
[1] 4898
```

```
> colnames(data6)
[1] "fixed.acidity"      "volatile.acidity"
[3] "citric.acid"        "residual.sugar"
[5] "chlorides"          "free.sulfur.dioxide"
[7] "total.sulfur.dioxide" "density"
[9] "pH"                 "sulphates"
[11] "alcohol"            "quality"
```



## Data Overview

The summary function is a built-in R function used to produce result summaries of various model fitting functions.

```
> summary(data6)
fixed.acidity    volatile.acidity    citric.acid
Min.      : 3.800      Min.      :0.0800      Min.      :0.0000
1st Qu.: 6.300      1st Qu.:0.2100      1st Qu.:0.2700
Median : 6.800      Median :0.2600      Median :0.3200
Mean      : 6.855      Mean      :0.2782      Mean      :0.3342
3rd Qu.: 7.300      3rd Qu.:0.3200      3rd Qu.:0.3900
Max.      :14.200      Max.      :1.1000      Max.      :1.6600
residual.sugar    chlorides    free.sulfur.dioxide
Min.      : 0.600      Min.      :0.00900      Min.      : 2.00
1st Qu.: 1.700      1st Qu.:0.03600      1st Qu.: 23.00
Median : 5.200      Median :0.04300      Median : 34.00
Mean      : 6.391      Mean      :0.04577      Mean      : 35.31
3rd Qu.: 9.900      3rd Qu.:0.05000      3rd Qu.: 46.00
Max.      :65.800      Max.      :0.34600      Max.      :289.00
total.sulfur.dioxide    density    pH
Min.      : 9.0      Min.      :0.9871      Min.      :2.720
1st Qu.:108.0      1st Qu.:0.9917      1st Qu.:3.090
Median :134.0      Median :0.9937      Median : 3.180
Mean      :138.4      Mean      :0.9940      Mean      : 3.188
3rd Qu.:167.0      3rd Qu.:0.9961      3rd Qu.:3.280
Max.      :440.0      Max.      :1.0390      Max.      :3.820
sulphates    alcohol    quality
Min.      :0.2200      Min.      : 8.00      Min.      :3.000
1st Qu.:0.4100      1st Qu.: 9.50      1st Qu.:5.000
Median :0.4700      Median :10.40      Median :6.000
Mean      :0.4898      Mean      :10.51      Mean      :5.878
3rd Qu.:0.5500      3rd Qu.:11.40      3rd Qu.:6.000
Max.      :1.0800      Max.      :14.20      Max.      :9.000
```


# Clustering



```
14 ▾ wssplot <- function(data, nc=15, seed=1234){  
15     wss <- (nrow(data)-1)*sum(apply(data,2,var))  
16 ▾   for (i in 2:nc){  
17       set.seed(seed)  
18 ▲     wss[i] <- sum(kmeans(data, centers=i)$withinss)}  
19     plot(1:nc, wss, type="b", xlab="Number of Clusters",  
20         ylab="Within groups sum of squares") +  
21       geom_vline(xintercept = 4, linetype = 1)  
22     wss  
23 ▲ }  
24  
25 wssplot(data7)  
26 KM = kmeans(data7,5)  
27 autoplot(KM,data7,frame=TRUE)  
28 #Cluster centers  
29 KM$centers
```




# Visualizations Clustering



# Visualizations Clustering

K means clusters for Alcohol vs. Fixed Acidity






# Key insights from Clustering model




## Elbow method

This method shows a clear view of how many clusters there should be.




## Clear clusters

There are 5 clear clusters in the plot and those refer to various different wines



## Balance for alcohol and acidity

Alcohol plays a big part in taste of wine, therefore a balance between alcohol and acidity plays a big part



## Easy of code

This method can be applied to other features as well for wine segmentation

# Linear Regression

```
# Importing dataset
dataset <- read.csv("winequality-white.csv")

# Correlation
cor(dataset, use = "everything")
cormat <- round(cor(dataset),2)
library(reshape2)
melted_cormat <- melt(cormat)

# Visualising the correlation
library(ggplot2)
ggplot(data = melted_cormat, aes(x=Var1, y=Var2, fill=value)) +
  geom_tile()
cona <- colnames(dataset)


# Splitting the dataset into the Training set and test set
library(caTools)
set.seed(123)
split = sample.split(dataset$alcohol, SplitRatio = 0.8)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)

# Fitting Multiple Linear Regression to the Training set
regressor = lm(formula = alcohol ~ .,
               data = training_set)
summary(regressor)
```

```
# Implementing automatic Backward Elimination of those variables not
# statistically significant
backwardElimination <- function(x, sl) {numVars = length(x)
  for (i in c(1:numVars)){regressor = lm(formula = alcohol ~ ., data = x)
    maxVar = max(coef(summary(regressor))[c(2:numVars), "Pr(>|t|)"])
    if (maxVar > sl){
      j = which(coef(summary(regressor))[c(2:numVars), "Pr(>|t|)"] == maxVar)
      x = x[, -j]
    }
    numVars = numVars - 1
  }
  return(summary(regressor))
}


SL = 0.05
dataset = dataset[, c(1,2,3,4,5,6,7,8,9,10,11,12)]
backwardElimination(training_set, SL)

# Predicting the Test set results
y_pred = predict(regressor, newdata = test_set)
```




# Visualizations Linear Regression

## Correlation Heatmap



## Alcohol vs. Fixed Acidity



# Linear Regression Results

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	7.132e+02	5.074e+00	140.576	< 2e-16	***
fixed.acidity	5.323e-01	9.081e-03	58.613	< 2e-16	***
volatile.acidity	5.945e-01	6.221e-02	9.555	< 2e-16	***
citric.acid	2.683e-01	5.219e-02	5.140	2.88e-07	***
residual.sugar	2.424e-01	2.680e-03	90.472	< 2e-16	***
free.sulfur.dioxide	-3.653e-03	4.582e-04	-7.974	2.00e-15	***
total.sulfur.dioxide	1.798e-03	2.105e-04	8.538	< 2e-16	***
density	-7.209e+02	5.180e+00	-139.178	< 2e-16	***
pH	2.379e+00	4.734e-02	50.260	< 2e-16	***
sulphates	9.792e-01	5.336e-02	18.351	< 2e-16	***
quality	3.956e-02	7.870e-03	5.027	5.20e-07	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3717 on 3912 degrees of freedom

Multiple R-squared: 0.9092, Adjusted R-squared: 0.9089

F-statistic: 3915 on 10 and 3912 DF, p-value: < 2.2e-16

# Key insights from linear regression model



## High accuracy

Model obtained an adjusted R-squared of 0.909




## Few Features

Implemented backward elimination to remove features not statistically significant



## Minimum code

Code can be implemented in quick succession and with all features. (Scores will vary)



## Feature correlation

Obtained a high R-squared even with many attributes with low correlation

# Logistic Regression

```
# Logistic Regression - Wine Quality

# Importing the dataset
dataset = read.csv('winequality-white.csv')

# Splitting the dataset into the Training set and Test set
# install.packages('caTools')
library(caTools)
set.seed(123)
dataset$quality = ifelse(dataset$quality/10 > 0.5, 1, 0)
split = sample.split(dataset$quality, SplitRatio = 0.80)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)

# Feature Scaling
training_set[-12] = scale(training_set[-12])
test_set[-12] = scale(test_set[-12])

# Fitting Logistic Regression to the Training set
classifier = glm(formula = quality ~ .,
                 family = binomial,
                 data = training_set)

# Predicting the Test set results
prob_pred = predict(classifier, type = 'response', newdata = test_set[-12])
y_pred = ifelse(prob_pred > 0.5, 1, 0)

# Making the Confusion Matrix
cm = table(ifelse(test_set[, 12]>0.5,1,0), y_pred)
print(cm)
print(sum(cm[1,1],cm[2,2])/sum(cm))
```

We used Logistic regression to predict the quality of white wine based on the chemical attributes of the wine. 0 is not good, 1 is good. we converted the data by using our quality score divided by 10. For those score above 0.5, we identified as good (1), below 0.5 as bad (0) .

Accuracy is 0.752.

```
> print(cm)
  y_pred
    0    1
0 157 171
1  72 580
> print(sum(cm[1,1],cm[2,2])/sum(cm))
[1] 0.7520408
```

# Naive Bayes

```
# Naive Bayes
# Importing the dataset
dataset = read.csv('winequality-white.csv')

# Splitting the dataset into the Training set and Test set
library(caTools)
set.seed(123)
dataset$quality = ifelse(dataset$quality/10 > 0.5, 1, 0)
split = sample.split(dataset$quality, SplitRatio = 0.80)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)

# Feature Scaling
training_set[,-12] = scale(training_set[,-12])
test_set[,-12] = scale(test_set[,-12])

# Fitting SVM to the Training set
# install.packages('e1071')
library(e1071)
classifier = naiveBayes(x = training_set[,-12],
                        y = training_set$quality)

# Predicting the Test set results
y_pred = predict(classifier, newdata = test_set[,-12])

# Making the Confusion Matrix
cm = table(test_set[, 12], y_pred)
print(cm)
print((sum(cm[1,1],cm[2,2])/sum(cm)))
```

We also tried Naive Bayes to predict the quality of white wine based on the chemical attributes of the wine. 0 is not good, 1 is good. we converted the data by using our quality score divided by 10. For those score above 0.5, we identified as good (1), below 0.5 as bad (0) .  
Accuracy: 0.6939.

```
> print(cm)
  y_pred
      0   1
0 179 149
1 151 501
> print(sum(cm[1,1],cm[2,2])/sum(cm))
[1] 0.6938776
```

# Summary

In this project, we used a white wine dataset to predict alcohol levels based on physical and chemical properties.

At the beginning of the project, we did a review of the data, deduced the main indicators of our data set.

We then applied the clustering method to look for patterns.

Then we tried 3 different models Linear, Logistic, Naive Bayes, the accuracy of the linear model is good, so we chose it as the best model.

For logistic regression, we created training and testing sets that we used to build the model. After we applied the Naive Bayes classifier, these algorithms resulted in a prediction accuracy of about 70 percent, but we were not satisfied with the results.



# References

[1] - Cortez, Paulo, Cerdeira, A., Almeida, F., Matos, T. & Reis, J.. (2009). Wine Quality. UCI Machine Learning Repository.